

[Return to Classroom](#)[DISCUSS ON STUDENT HUB](#)

# Collaboration and Competition

REVIEW

CODE REVIEW 2

HISTORY

Meets Specifications

Congratulations on passing your final project 🎉🎊

Assuming you passed the first two projects, THIS IS GRADUATION

You made a great submission that solved this quite unstable environment

Rate this review

START

episodes. It's common for this environment to give variant results every run, but your model is stable. Give yourself proper credit.

---

To answer your student note, Batch Normalization was introduced in a 2015 [paper](#) to address the vanishing/exploding gradients problems. This operation simply zero-centers and normalizes each input, then scales and shifts the result using two new parameter vectors per layer: one for scaling, one for shifting. Basically, Batch Normalization lets the model learn the optimal scale and mean of each of the layer's inputs. Thus speeding the learning process.

THIS IS A GREAT QUESTION. I SUGGEST YOU READ THE PAPER CAREFULLY AS THIS IS A VERY COMMON AND BENEFICIAL TECHNIQUE.

---

I hope this nanodegree gives you a good kick towards your dreams. These resources will give you a different point of view for Deep Reinforcement Learning:

- [MIT 6.S094: Deep Reinforcement Learning video.](#)
- [Deep Reinforcement Learning by David Silver.](#)

## Training Code

The repository includes functional, well-documented, and organized code for training the agent.

Your code functions correctly. Here's the result I got when I ran it:

Rate this review

START

Episode 100	Average Score: 0.010	Moving Avg: 0.010
Episode 200	Average Score: 0.011	Moving Avg: 0.011
Episode 300	Average Score: 0.014	Moving Avg: 0.014
Episode 400	Average Score: 0.020	Moving Avg: 0.020
Episode 500	Average Score: 0.039	Moving Avg: 0.029
Episode 600	Average Score: 0.048	Moving Avg: 0.038
Episode 700	Average Score: 0.064	Moving Avg: 0.064
Episode 800	Average Score: 0.108	Moving Avg: 0.098
Episode 900	Average Score: 0.191	Moving Avg: 0.191
Episode 1000	Average Score: 0.318	Moving Avg: 0.308
Episode 1028	Average Score: 0.326	Moving Avg: 0.326

*PS: I didn't have enough time to wait for the training to finish.*

As you see, the results are similar to yours:

Episode 100	Average Score: 0.000	Moving Avg: 0.000
Episode 200	Average Score: 0.001	Moving Avg: 0.001
Episode 300	Average Score: 0.003	Moving Avg: 0.003
Episode 400	Average Score: 0.002	Moving Avg: 0.002
Episode 500	Average Score: 0.003	Moving Avg: 0.003
Episode 600	Average Score: 0.016	Moving Avg: 0.006
Episode 700	Average Score: 0.013	Moving Avg: 0.013
Episode 800	Average Score: 0.029	Moving Avg: 0.019
Episode 900	Average Score: 0.023	Moving Avg: 0.023
Episode 1000	Average Score: 0.034	Moving Avg: 0.034
Episode 1100	Average Score: 0.076	Moving Avg: 0.066
Episode 1200	Average Score: 0.119	Moving Avg: 0.109
Episode 1300	Average Score: 0.278	Moving Avg: 0.268
Episode 1400	Average Score: 0.439	Moving Avg: 0.429
Episode 1467	Average Score: 0.500	Moving Avg: 0.500
Environment solved in 1467 episodes!		Average Score: 0.50

Rate this review

START

The code is written in PyTorch and Python 3.

The submission includes the saved model weights of the successful agent.

The model weights of the successful agent are saved in `checkpoint_final.pth`.

## README

The GitHub submission includes a `README.md` file in the root of the repository.

The README describes the the project environment details (i.e., the state and action spaces, and when the environment is considered solved).

The `Introduction` of the README **correctly** describes all the project environment details.

The README has instructions for installing dependencies or downloading needed files.

The `Getting Started` section of the README gives correct instructions for:

- ✓ Downloading the environment.
- ✓ Installing dependencies.

The README describes how to run the code in the repository, to train the agent. For additional resources on creating READMEs or using Markdown, see [here](#) and [here](#).

The `Instructions` section of the README **correctly** describes how to run your code.

Rate this review

START

## Report

The submission includes a file in the root of the GitHub repository (one of `Report.md`, `Report.ipynb`, or `Report.pdf`) that provides a description of the implementation.

The report clearly describes the learning algorithm, along with the chosen hyperparameters. It also describes the model architectures for any neural networks.

### All requirements are explained.

- ✓ Hyperparameters in the `Implementation` section.
- ✓ Learning Algorithm in the `MADDPG – Agent` section.
- ☐ Model Architecture in the `Implementation` section. Please mention using `tanh` on one of the layers.

YOU ALREADY PASSED THIS POINT. YET, IT'S MISSING THIS ONE DETAIL. PART OF OUR JOB IS TO GUIDE YOU IN BUILDING A GREAT PORTFOLIO. I HOPE THIS HELPS.

A plot of rewards per episode is included to illustrate that the agents get an average score of +0.5 (over 100 consecutive episodes, after taking the maximum over both agents).

The submission reports the number of episodes needed to solve the environment.

Great job 🙌 Your agents train in a reasonable number of episodes demonstrated in the `Implementation` section.

- ✓ The Plot for the trained agents is included.
- ✓ The Report mentions the number of episodes needed to solve the environment.

Rate this review

START

The submission has concrete future ideas for improving the agent's performance.

- Yes, you can almost always fine-tune the hyperparameters and the model. Grid Search is a very good suggestion.
- Have fun with the football version using shared communications.

You can also try using different algorithms. Here are resources to help:

- [A3C](#)
- [PPO](#) and a [PPO implementation](#).
- [D4PG](#)

 [DOWNLOAD PROJECT](#)

2

[CODE REVIEW COMMENTS](#)



Rate this review

START

RETURN TO PATH

Rate this review

START